

Ex. 1.12.1: Conditional probability: suppose that if $\theta = 1$, then y has a normal distribution with mean 1 and standard deviation σ , and if $\theta = 2$, then y has a normal distribution with mean 2 and standard deviation σ . Also, suppose $\Pr(\theta = 1) = 0.5$ and $\Pr(\theta = 2) = 0.5$.

- For $\theta = 2$, write the formula for the marginal probability density for y and sketch it.
- What is $\Pr(\theta = 1|y = 1)$, again supposing $\sigma = 2$?
- Describe how the posterior density of θ changes in shape as σ is increased and as it is decreased.

Answer:

The formula for the marginal probability is that of $\mathcal{N}(\theta, \sigma^2)$:

$$p(y|\theta = 2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2}\left(\frac{y-2}{\sigma}\right)^2\right).$$

As for the sketch, I'm too lazy to do it. It's a bell shape, centered at 2 and σ being half the width at about 0.61 height. There's an xkcd-looking Matplotlib plot in Figure 1.




Figure 1: Your typical bell thingy with $\sigma = 2$ and mean 2

Now then,

$$\begin{aligned} \Pr(\theta = 1|y = 1) &= \frac{\Pr(\theta = 1) \Pr(y = 1|\theta = 1)}{\Pr(\theta = 1) \Pr(y = 1|\theta = 1) + \Pr(\theta = 2) \Pr(y = 1|\theta = 2)} \\ &= \frac{0.5 \cdot \frac{1}{\sqrt{8\pi}}}{0.5 \cdot \frac{1}{\sqrt{8\pi}} + 0.5 \cdot \frac{1}{\sqrt{8\pi}} \exp\left(-\frac{1}{8}\right)} \\ &\approx 0.53. \end{aligned}$$

The posterior for θ gets more homogeneous as σ increases, and instead gets closer to $(1, 0)$ otherwise, since σ directly defines the overlap between both likelihood functions as functions of y .

Ex. 1.12.2: Conditional means and variances: show that (1.8) and (1.9) hold if u is a vector.

Answer:

No, u are a vector.

Just kidding. Proof of (1.8) is symbolically the same. As to (1.9), it is *about* the same. First, note that

$$\begin{aligned} \text{var}(x) &= \mathbb{E} \left[(x - \mathbb{E}(x)) (x - \mathbb{E}(x))^T \right] \\ &= \mathbb{E} \left[(x - \mathbb{E}(x)) \left(x^T - (\mathbb{E}(x))^T \right) \right] \\ &= \mathbb{E} \left[xx^T - x (\mathbb{E}(x))^T - \mathbb{E}(x) x^T + (\mathbb{E}(x)) (\mathbb{E}(x))^T \right] \\ &= \mathbb{E} (xx^T) - (\mathbb{E}(x)) (\mathbb{E}(x))^T. \end{aligned}$$

The result then follows replacing any occurrence of x^2 by xx^T .

Ex. 1.12.3: Probability calculation for genetics (from Lindley, 1965): suppose that in each individual of a large population there is a pair of genes, each of which can be either x or X , that controls eye color: those with xx have blue eyes, while heterozygotes (those with Xx or xX) and those with XX have brown eyes. The proportion of blue-eyed individuals is p^2 and of heterozygotes is $2p(1 - p)$, where $0 < p < 1$. Each parent transmits one of its own genes to the child; if a parent is a heterozygote, the probability that it transmits the gene of type X is $\frac{1}{2}$. Assuming random mating, show that among brown-eyed children of brown-eyed parents, the expected proportion of heterozygotes is $2p/(1 + 2p)$. Suppose Judy, a brown-eyed child of brown-eyed parents, marries a heterozygote, and they have n children, all brown-eyed. Find the posterior probability that Judy is a heterozygote and the probability that her first grandchild has blue eyes.

Answer:

Denoting p_1 and p_2 both parents' alleles, random mating implies $p(p_1|p_2)$ and $p(p_2)$ are both given by the population's proportion of people with p_2 and p_1 after p_2 . We'll also use the hypothesis of a large pool of people, so that p remains approximately equal after picking one of the parents. We thus have $p(p_1, p_2) = p(p_1)p(p_2)$.

So, if both parents are brown-eyed, they're one from $\{xX, Xx, XX\}$. The conditional probability of being any of the first two is

$$\begin{aligned} \Pr(xX + Xx | xX + Xx + XX) &= \frac{\Pr((xX + Xx)(xX + Xx + XX))}{\Pr(xX + Xx + XX)} \\ &= \frac{2p(p - 1)}{2p(p - 1) + (1 - p)^2} \\ &= \frac{2p}{1 + p}. \end{aligned}$$

And the probability of transmitting either allele is then $1/2$. The probability of being an heterozygote is twice the probability that one of the parent gives x and the other X , so

$$\begin{aligned} &2 \left(\frac{1}{2} \frac{2p}{1 + p} \right) \left(\frac{1}{2} \frac{2p}{1 + p} + \frac{1 - p}{1 + p} \right) \\ &2 \left(\frac{1}{2} \frac{2p}{1 + p} \right) \left(\frac{1}{1 + p} \right) \\ &\frac{2p}{(1 + p)^2}. \end{aligned}$$

Finally, since we want this conditioned to the fact that they're all brown eyed, we must quotient this with the probability of being brown eyed, which is the above quantity plus the probability that both parents give X , $1/(1 + p)^2$, so the final proportion is

$$\begin{aligned} &\left[\frac{2p}{(1 + p)^2} \right] \left[\frac{2p}{(1 + p)^2} + \frac{1}{(1 + p)^2} \right]^{-1} \\ &= \frac{2p}{2p + 1}. \end{aligned}$$

We are now asked the posterior probability of Judy being heterozygote after blah blah blah. So the priors of them being and not being heterozygote are $2p/(2p + 1)$ and $1/(2p + 1)$. The likelihood of the n childrens are $\left(\frac{3}{4}\right)^n$ and 1 respectively, so that the posterior probability is

$$\frac{2p}{2p + \left(\frac{4}{3}\right)^n},$$

which of course tends to 0 as n increases, since the likelihood of not having a blue eyed child decreases as n increases.

Ex. 1.12.4: Probability assignment: we will use the football dataset to estimate some conditional probabilities about professional football games. There were twelve games with point spread of 8 points; the outcomes in those

games were: $-7, -5, -3, -3, 1, 6, 7, 13, 15, 16, 20$, and 21 , with positive values indicating wins by the favorite and negative values indicating wins by the underdog. Consider the following conditional probabilities:

$$\begin{aligned}\Pr(\text{favorite wins}|\text{point spread} = 8), \\ \Pr(\text{favorite wins by at least 8}|\text{point spread} = 8), \\ \Pr(\text{favorite wins by at least 8}|\text{point spread} = 8 \text{ and favorite wins}).\end{aligned}$$

- a Estimate each of these using the relative frequencies of games with a point spread of 8.
- b Estimate each using the normal approximation for the distribution of (outcome $-$ point spread).

Answer:

Point a is just counting:

$$\begin{aligned}\Pr(\text{favorite wins}|\text{point spread} = 8) &= 8/12 \approx 0.58 \\ \Pr(\text{favorite wins by at least 8}|\text{point spread} = 8) &= 5/12 \approx 0.42, \\ \Pr(\text{favorite wins by at least 8}|\text{point spread} = 8 \text{ and favorite wins}) &= 5/8 \approx 0.63.\end{aligned}$$

Point b uses the model $P(\text{wins by } \hat{y}|\text{point spread of } x) = \mathcal{N}(\hat{y}|x, 14^2)$, so

$$\begin{aligned}\Pr(\text{favorite wins}|\text{point spread} = 8) &= 1 - \Phi\left(\frac{0 - 8}{14}\right) \\ &\approx 0.72 \\ \Pr(\text{favorite wins by at least 8}|\text{point spread} = 8) &= 0.5, \\ \Pr(\text{favorite wins by at least 8}|\text{point spread} = 8 \text{ and favorite wins}) &\approx \frac{0.5}{0.72} \\ &\approx 0.69\end{aligned}$$

Ex. 1.12.5: Probability assignment: the 435 U.S. Congressmembers are elected to two-year terms; the number of voters in an individual congressional election varies from 50 000 to 350 000. We will use various sources of information to estimate roughly the probability that at least one congressional election is tied in the next national election.

- a Use any knowledge you have about U.S. politics. Specify clearly what information you are using to construct this conditional probability, even if your answer is just a guess.
- b Use the following information: in the period 1900 $-$ 1992, there were 20 597 congressional elections, out of which 6 were decided by fewer than 10 votes and 49 by fewer than 100 votes.

Answer:

Got no knowledge. Gonna skip this one.

Ex. 1.12.6: Conditional probability: approximately $1/125$ of all births are fraternal twins and $1/300$ of births are identical twins. Elvis Presley had a twin brother (who died at birth). What is the probability that Elvis was an identical twin? (You may approximate the probability of a boy or girl birth as $1/2$).

Answer:

We need

$$P(\text{Identical}|\text{Both male}) = \frac{P(\text{Both male}|\text{Identical})P(\text{Identical})}{P(\text{Both male}|\text{Identical})P(\text{Identical}) + P(\text{Both male}|\text{Fraternal})P(\text{Fraternal})}.$$

This is pretty straightforward:

$$\begin{aligned} P(\text{Both male}|\text{Identical}) &= \frac{1}{2}, \\ P(\text{Identical}) &= \frac{1}{300}, \\ P(\text{Both male}|\text{Fraternal}) &= \frac{1}{2} \frac{1}{2}, \\ P(\text{Fraternal}) &= \frac{1}{125} \end{aligned}$$

It then follows that

$$P(\text{Identical}|\text{Both male}) = 5/11 \approx 0.45.$$

Ex. 1.12.7: Conditional probability: the following problem is loosely based on the television game show *Let's Make a Deal*. At the end of the show, a contestant is asked to choose one of three large boxes, where one box contains a fabulous prize and the other two boxes contain lesser prizes. After the contestant chooses a box, Monty Hall, the host of the show, opens one of the two boxes containing smaller prizes. (In order to keep the conclusion suspenseful, Monty does not open the box selected by the contestant.) Monty offers the contestant the opportunity to switch from the chosen box to the remaining unopened box. Should the contestant switch or stay with the original choice? Calculate the probability that the contestant wins under each strategy. This is an exercise in being clear about the information that should be conditioned on when constructing a probability judgement.

Answer:

Let W_i = the winning box is the i th, P_i = the player chose the i th box, H_i = the host opened the i th box. We then want

$$\begin{aligned} P(W_1|P_2H_3) &= \frac{P(H_3|P_2W_1)P(W_1|P_2)}{P(H_3|P_2W_1)P(W_1|P_2) + P(H_3|P_2W_2)P(W_2|P_2) + P(H_3|P_2W_3)P(W_1|P_3)} \\ &= \frac{1}{1 + \frac{1}{2} + 0} \\ &= \frac{2}{3}. \end{aligned}$$

So, the obvious choice is to switch.

Ex. 1.12.8: Subjective probability: discuss the following statement. ‘The probability of event E is considered “subjective” if two rational persons A and B can assign unequal probabilities to E , $P_A(E)$ and $P_B(E)$. These probabilities can also be interpreted as “conditional”: $P_A(E) = P(E|I_A)$ and $P_B(E) = P(E|I_B)$, where I_A and I_B represent the knowledge available to persons A and B , respectively’. Apply this idea to the following examples.

- a The probability that a ‘6’ appears when a fair die is rolled, where A observes the outcome of the die roll and B does not.
- b The probability that Brazil wins the next World Cup, where A is ignorant of soccer and B is a knowledgeable sports fan.

Answer:

I’d say both are subjective. The knowledge of the die throw *before* observation is that there’s 1/6 probability for any side, since it’s a fair die. But if A knows the result of the die, then the probability distribution should no longer incorporate any uncertainty. The second example is more evident, since it is expected that people more knowledgeable have a better understanding of their degree of knowledge.

Ex. 1.12.9: Simulation of a queuing problem: a clinic has three doctors. Patients come into the clinic at random, starting at 9 a.m., according to a Poisson process with time parameter 10 minutes: that is, the time after opening at which the first patient appears follows an exponential distribution with expectation 10 minutes and then, after each patient arrives, the waiting time until the next patient is independently exponentially distributed, also with expectation 10 minutes. When a patient arrives, he or she waits until a doctor is available. The amount of time

spent by each doctor with each patient is a random variable, uniformly distributed between 5 and 20 minutes. The office stops admitting new patients at 4 p.m. and closes when the last patient is through with the doctor.

- Simulate this process once. How many patients came to the office? How many had to wait for a doctor? What was their average wait? When did the office close?
- Simulate the process 100 times and estimate the median and 50% interval for each of the summaries in (a).

Answer:

Script “1.12.9.py” generates these answers:

48 patients came to the office.

38 patients had to wait for a doctor.

Average wait time was 14.09 minutes.

Office closed at 16 : 02 : 00.

The rest of answers can be found in Figure 2.

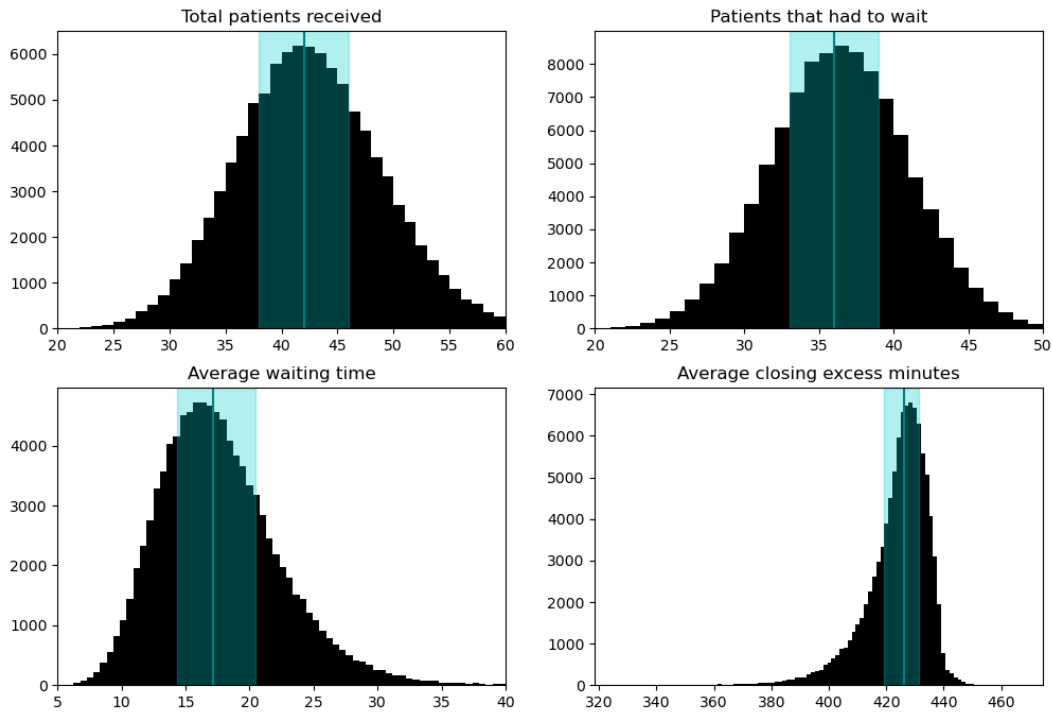


Figure 2: Histograms for the stuff after 100000 runs.

Ex. 2.11.1: Posterior inference: suppose you have a $\text{Beta}(4, 4)$ prior distribution on the probability θ that a coin will yield a ‘head’ when spun in a specified manner. The coin is independently spun ten times, and ‘heads’ appear fewer than 3 times. You are not told how many heads were seen, only that the number is less than 3. Calculate your exact posterior density (up to a proportionality constant) for θ and sketch it.

Answer:

We are asked

$$\begin{aligned}
 p(\theta|y < 3) &= p(\theta|(y = 0) + (y = 1) + (y = 2)) \\
 &\propto p((y = 0) + (y = 1) + (y = 2)|\theta)p(\theta) \\
 &= (p(y = 0|\theta) + p(y = 1|\theta) + p(y = 2|\theta))p(\theta).
 \end{aligned}$$

Each term separately will yield $\text{Beta}(4 + y, 4 + (10 - y))$, so that the final probability distribution is plainly

$$\frac{1}{3} [\text{Beta}(4, 14) + \text{Beta}(5, 13) + \text{Beta}(6, 12)].$$

Those have means $4/18$, $5/18$, $6/18$, so that the mean is $5/18$. More concretely, if

$$p(x) = \sum_i w_i p_i(x),$$

then the expectation of stuff is given by

$$\begin{aligned} E \cdot &= \int dx \, p(x) \cdot \\ &= \sum_i w_i \int dx \, p_i(x) \cdot \\ &= \sum_i w_i E_i \cdot. \end{aligned}$$

It then follows that

$$\begin{aligned} E\theta &= \frac{1}{3} [E_1\theta + E_2\theta + E_3\theta] \\ &= \frac{5}{18}. \end{aligned}$$

The same process can yield the new variance.

The corresponding computer graph is shown in Figure 3.




Figure 3: Probability distribution for θ given the number of heads is less than 3.

Ex. 2.11.2: Predictive distributions: consider two coins, C_1 and C_2 , with the following characteristics: $\Pr(\text{heads}|C_1) = 0.6$ and $\Pr(\text{heads}|C_2) = 0.4$. Choose one of the coins at random and imagine spinning it repeatedly. Given that the first two spins from the chosen coin are tails, what is the expectation of the number of additional spins until a head shows up?

Answer:

Here I'll represent the data by D .

The conditional probability that we make n tosses until heads is found is

$$p(n|C_i) = (\Pr(\text{tails}|C_i))^{n-1} \Pr(\text{heads}|C_i),$$

so that

$$\begin{aligned} p(n|D) &= \sum_i p(n, C_i|D) \\ &= \sum_i p(n|C_i) \Pr(C_i|D). \end{aligned}$$

Since we have no information as to which coin we picked, we'll take $\Pr(C_i) = \frac{1}{2}$ which, together with $\Pr(D|C_i) = (\Pr(\text{tails}|C_i))^2$ yield $\Pr(C_i|D) = (\Pr(\text{tails}|C_i))^2 / [(\Pr(\text{tails}|C_i))^2 + (\Pr(\text{tails}|C_{\bar{i}}))^2]$, where \bar{i} is not i . Numerically, this yields $P(C_2|D) = 0.69$, nice.

Anyways, we also have $p(n|C_i) = (\Pr(\text{tails}|C_i))^{n-1} \Pr(\text{heads}|C_i)$, so we get our final result of

$$p(n|D) = \frac{(\Pr(\text{tails}|C_1))^{n+1} \Pr(\text{heads}|C_1) + (\Pr(\text{tails}|C_2))^{n+1} \Pr(\text{heads}|C_2)}{(\Pr(\text{tails}|C_1))^2 + (\Pr(\text{tails}|C_2))^2}.$$

Our final calculation is then $E[n] = \sum_{n=1} np(n)$, which can be done analitically just fine but I wont, and yields about 2.24.

Ex. 2.11.3: Predictive distributions: let y be the number of 6's in 1000 rolls of a fair die.

- Sketch the approximate distribution of y , based on the normal approximation.
- Using the normal distribution table, give approximate 5%, 25%, 50%, 75% and 95% points for the distribution of y .

Answer:

We have $y \sim B(1000, 1/6)$, with mean $166.\bar{6}$ and deviation 11.8. So I should draw a bell curve which should have a width of 23.6 at 0.6 its height, centered in 166.6. As for the table thingy, the following Python interpreter run has been performed

```
In [1]: import scipy.stats as st

In [2]: [ 166.6+11.8*st.norm.ppf(x) for x in [0.05, 1/4, 1/2, 3/4, 0.95] ]
Out[2]:
[147.1907272019726,
 158.64102094768623,
 166.6,
 174.55897905231376,
 186.00927279802738]
```

Ex. 2.11.4: Predictive distributions: let y be the number of 6's in 1000 independent rolls of a particular real die, which may be unfair. Let θ be the probability that the die lands on '6'.

Suppose your prior distribution for θ is as follows:

$$\begin{aligned} \Pr(\theta = 1/12) &= 0.25 \\ \Pr(\theta = 1/6) &= 0.5 \\ \Pr(\theta = 1/4) &= 0.25. \end{aligned}$$

- Using the normal approximation for the conditional distributions, $p(y|\theta)$, sketch your approximate prior predictive distribution for y .
- Give approximate 5%, 25%, 50%, 75%, and 95% points for the distribution of y . (Be careful here: y does not have a normal distribution, but you can still use the normal distribution as part of your analysis.)

Answer:

Again we have $p(y) = \sum_{\theta} p(y|\theta)p(\theta)$. Under the normal approximation, $y|\theta \sim \mathcal{N}(n\theta, n\theta(1-\theta))$, so we'll have

$$y|\theta \sim 0.25 \cdot N(83.3, 8.74^2) + 0.5 \cdot N(166.7, 11.8^2) + 0.25 \cdot N(250, 13.7),$$




Figure 4: Actual and normally-approximated predictive prior distribution for y

which is a pretty disjoint (and thus easily drawable) mixture if you ask me.

As for the approximate points, we can essentially regard the distributions as disjoint, so that the 5% point is the 20% of the first one, the 25% can be approximated as, say, the middle point between the first two Gaussians, the 50% point as the mean 166.7 of the middle Gaussian, and so on. So, *a ojo*, they should be about 70, 120, 167, 210, and 270. But I shall put this handwavery to the test via numerical integration, which is performed in the script ‘2.11.4.py’. For the real values it yields 76, 120, 167, 207, and 261. For the normal-plus-disjointness-approximated values we get 76, 125, 167, 208, and 262.

Ex. 2.11.5: Posterior distribution as a compromise between prior information and data: let y be the number of heads in n spins of a coin, whose probability of heads is θ .

- a If your prior distribution for θ is uniform on the range $[0, 1]$, derive your prior predictive distribution for y ,

$$\Pr(y = k) = \int_0^1 \Pr(y = k|\theta)d\theta,$$

for each $k = 0, 1, \dots, n$.

- b Suppose you assign a $\text{Beta}(\alpha, \beta)$ prior distribution for θ , and then you observe y heads out of n spins. Show algebraically that your posterior mean of θ always lies between your prior mean, $\frac{\alpha}{\alpha+\beta}$, and the observed relative frequency of heads, $\frac{y}{n}$.
- c Show that, if the prior distribution on θ is uniform, the posterior variance of θ is always less than the prior variance.
- d Give an example of a $\text{Beta}(\alpha, \beta)$ prior distribution and data y, n , in which the posterior variance of θ is higher than the prior variance.

Answer:

Point a:

$$\begin{aligned}
\Pr(y = k) &= \int_0^1 d\theta \Pr(y = k|\theta) \\
&= \int_0^1 d\theta \binom{n}{k} \theta^k (1 - \theta)^{n-k} \\
&= \binom{n}{k} \int_0^1 d\theta \theta^k (1 - \theta)^{n-k} \\
&= \binom{n}{k} \left\{ \left[\frac{1}{k+1} \theta^{k+1} (1 - \theta)^{n-k} \right]_0^1 + \frac{n-k}{k+1} \int_0^1 d\theta \theta^{k+1} (1 - \theta)^{n-k-1} \right\} \\
&= \binom{n}{k} \frac{n-k}{k+1} \frac{n-k-1}{k+2} \cdots \frac{1}{n} \int_0^1 d\theta \theta^n \\
&= \binom{n}{k} \frac{k!(n-k)!}{(n+1)!} \\
&= \frac{1}{n+1}.
\end{aligned}$$

Point b:

We've seen in the book that we now have $\theta|y \sim \text{Beta}(y + \alpha, n - y + \beta)$, which has expectation value

$$\begin{aligned}
\frac{y + \alpha}{n + \alpha + \beta} &= \frac{y}{n} \frac{n}{n + \alpha + \beta} + \frac{\alpha}{\alpha + \beta} \frac{\alpha + \beta}{n + \alpha + \beta} \\
&= w \frac{y}{n} + (1 - w) \frac{\alpha}{\alpha + \beta} \quad \left(w = \frac{n}{n + \alpha + \beta} \in [0, 1] \right).
\end{aligned}$$

This is a linear function of w which goes from $\frac{y}{n}$ to $\frac{\alpha}{\alpha + \beta}$ as w goes from 0 to 1, thus proving the result.

Point c:

Since $\theta \sim 1$, its variance is

$$\begin{aligned}
\int_0^1 d\theta \left(\theta - \frac{1}{2} \right)^2 &= \frac{1}{3} \left[\left(\theta - \frac{1}{2} \right)^3 \right]_0^1 \\
&= \frac{1}{12}.
\end{aligned}$$

Now, $\theta|y \sim \text{Beta}(y + 1, n - y + 1)$, which has variance

$$\frac{(y + 1)(n - y + 1)}{(n + 2)^2(n + 3)}.$$

As a function of y , this is proportional to the denominator $-y^2 + ny + n + 1$, whose maximum is at $\frac{1}{2}n$ and yields $\frac{1}{4}(2n + 1)$, so that

$$\frac{(y + 1)(n - y + 1)}{(n + 2)^2(n + 3)} \leq \frac{1}{4} \frac{(2n + 1)^2}{(n + 2)^2(n + 3)}.$$

This function can be maximized analitically via derivation, but I just plotted it. It's maximum is around 3, taking the value $49/600 < 50/600 = 1/12$. Such a close call though.

As for point c, take $\text{Beta}(3, 1)$, yielding variance 0.0375, and suppose one negative outcome is measured, taking the posterior distribution towards $\text{Beta}(3, 2)$ with variance 0.04.

Ex. 2.11.6: Predictive distributions: Derive the mean and variance (2.17) of the negative binomial predictive distribution for the cancer rate example, using the mean and variance formulas (1.8) and (1.9).

Answer:

Formulas (1.8) and (1.9) are

$$\begin{aligned} E(u) &= E(E(u|v)), \\ \text{var}(u) &= E(\text{var}(u|v)) + \text{var}(E(u|v)). \end{aligned}$$

We're thus asked the mean and variance of the distribution $p(y) = \int d\theta p(y|\theta)p(\theta)$, where $\theta \sim \text{Gamma}(\alpha, \beta)$ and $y|\theta \sim \text{Poisson}(10n\theta)$. Direct application follows:

$$\begin{aligned} E(y) &= E(E(y|\theta)) \\ &= E(10n\theta) \\ &= 10n_j E(\theta) \\ &= 10n \frac{\alpha}{\beta}. \\ \text{var}(y) &= E(\text{var}(y|\theta)) + \text{var}(E(y|\theta)) \\ &= E(10n\theta) + \text{var}(10n\theta) \\ &= 10n \frac{\alpha}{\beta} + (10n)^2 \frac{\alpha}{\beta^2}, \end{aligned}$$

where the mean and variance of the Gamma and Poisson have been Googled and used straight away.

Ex. 2.11.7: Noninformative prior densities:

- a For the binomial likelihood, $y \sim \text{Bin}(n, \theta)$, show that $p(\theta) \propto \theta^{-1}(1 - \theta)^{-1}$ is the uniform prior distribution for the natural parameter of the exponential family.
- b Show that if $y = 0$ or n , the resulting posterior distribution is improper.

Answer:

We can express the binomial likelihood in exponential form as follows

$$\begin{aligned} \text{Bin}(y|n, \theta) &= \binom{n}{y} \theta^y (1 - \theta)^{n-y} \\ &= \binom{n}{y} (1 - \theta)^n \left(\frac{\theta}{1 - \theta} \right)^y \\ &= \binom{n}{y} (1 - \theta)^n \exp \left(y \log \left(\frac{\theta}{1 - \theta} \right) \right). \end{aligned}$$

It then becomes clear that the natural parameter is $\phi = \log \left(\frac{\theta}{1 - \theta} \right)$. Since the transformation between ϕ and θ is one-to-one,

$$\begin{aligned} p(\phi) &= p(\theta) \left| \frac{d\phi}{d\theta} \right|^{-1} \\ &= \theta^{-1} (1 - \theta)^{-1} \left| \frac{\theta}{1 - \theta} (1 - \theta)^2 \right| \\ &= 1. \end{aligned}$$

As for the posterior distribution, we have

$$\begin{aligned} p(\theta|y) &\propto p(y|\theta)p(\theta) \\ &\propto \theta^{y-1} (1 - \theta)^{n-y-1}. \end{aligned}$$

Whether $y = 0$ or $y = n$ we have something of the form $\eta^{-1}(1-\eta)^{n-1}$, η being either θ or $1-\theta$, so that the integral of this quantity over η is precisely the normalization required, but it yields

$$\begin{aligned} \int_0^1 d\eta \eta^{-1}(1-\eta)^{n-1} &= \int_0^1 d\eta \sum_{k=0}^{n-1} \binom{n-1}{k} \eta^{-1} \eta^k \\ &= \sum_{k=0}^{n-1} \binom{n-1}{k} \int_0^1 d\eta \eta^{k-1} \\ &= \int_0^1 d\eta \eta^{-1} + \sum_{k=1}^{n-1} \binom{n-1}{k} \int_0^1 d\eta \eta^{k-1} \\ &= \infty + \text{something finite}, \end{aligned}$$

whence this integral is improper.

Ex. 2.11.8: Normal distribution with unknown mean: a random sample of n students is drawn from a large population, and their weights are measured. The average weight of the n sampled students is $\bar{y} = 150$ pounds. Assume the weights in the population are normally distributed with unknown mean θ and known standard deviation 20 pounds. Suppose your prior distribution for θ is normal with mean 180 and standard deviation 40.

- Give your posterior distribution for θ . (Your answer will be a function of n .)
- A new student is sampled at random from the same population and has a weight of \tilde{y} pounds. Give a posterior predictive distribution for \tilde{y} . (Your answer will still be a function of n .)
- For $n = 10$, give a 95% posterior interval for θ and a 95% posterior predictive interval for \tilde{y} .
- Do the same for $n = 100$.

Answer:

By hypothesis, $y|\theta \sim N(\theta, 20^2)$ and $\theta \sim N(180, 40^2)$, so that $\bar{y}|\theta \sim N(\theta, 20^2/n)$, and thus $\theta|\bar{y}$ is gonna be given by a normal distribution with

$$\begin{aligned} E(\theta|\bar{y}) &= \frac{40^{-2}180 + n20^{-2}150}{40^{-2} + n20^{-2}}, \\ \text{var}(\theta|\bar{y}) &= \frac{1}{40^{-2} + n20^{-2}}. \end{aligned}$$

I could put some numbers there, but it'd be no use. I instead plotted probability density for $\theta|\bar{y}$ (vertical axis) as a function of sample size n (horizontal axis) for n between 0 and 100, as shown in Figure 5.




Figure 5: Non-normalized probability distribution for $\theta|\bar{y}$ as a function of θ and number of samples n . Black contour is the mean, and grey contours are 1 and 3 standard deviations from it.

Since the actual precision of the data is so small, the mean converges quickly towards the observed mean.

The posterior predictive distribution is the same, but incorporates an extra 20 in the standard deviation. In particular, for $n = 10$, the posterior probabilities for θ and \tilde{y} have mean 150.73 and respectively a standard deviation of 6.24 and 26.24, so that the $\sim 95\%$ intervals are 150.73 ± 12.48 and 150.73 ± 42.48 . For $n = 100$, the mean becomes 150.07 and the standard deviations 1.99 and 20.199, so that the intervals are 150.07 ± 4 and 150.07 ± 48 .

Ex. 2.11.9: Setting parameters for a beta prior distribution: suppose your prior distribution for θ , the proportion of Californians who support the death penalty is beta with mean 0.6 and standard deviation 0.3.

- Determine the parameters α and β of your prior distribution. Sketch the prior density function.
- A random sample of 1000 Californians is taken, and 65% support the death penalty. What are your posterior mean and variance for θ ? Draw the posterior density function.
- Examine the sensitivity of the posterior distribution to different prior means and widths including a non-informative prior.

Answer:

We equate the mean and the variance

$$\frac{\alpha}{\alpha + \beta} = \mu,$$

$$\frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)} = \sigma^2.$$

Replacing $\alpha/(\alpha + \beta) = \mu$ and $\beta/(\alpha + \beta) = 1 - \mu$ in the second equation yields

$$\alpha + \beta = \frac{\mu(1 - \mu)}{\sigma^2} - 1,$$

which into the first equation implies

$$\alpha = \mu \left(\frac{\mu(1 - \mu)}{\sigma^2} - 1 \right),$$

$$\beta = (1 - \mu) \left(\frac{\mu(1 - \mu)}{\sigma^2} - 1 \right).$$

Upon replacing the particular values $\mu = 0.6$ and $\sigma = 0.3$, we get

$$\alpha = 1, \quad \beta = \frac{2}{3}.$$

Now, after the census we get 650 positives and 350 negatives, whence we'll have $\theta|D \sim \text{Beta}(650 + 1, 350 + \frac{2}{3})$, with mean 65.0% and standard deviation 1.4%.

The following table, calculated in the corresponding script to this exercise, evaluates a few alternatives for σ , ranging from *whoa* prior certainty to complete prior uncertainty.

Prior deviation	0.1%	0.5%	1.0%	5.0%	10.0%	50.0%	100%	Noninformative
Posterior mean	60.02%	60.47%	61.47%	64.57%	64.89%	65.00%	65.00%	64.97%
Posterior deviation	0.10%	0.47%	0.83%	1.44%	1.49%	1.51%	1.51%	1.51%

The noninformative alternative diminishes the posterior mean because it's centered around 50% instead of 60%, so that it's "pulling effect" is more evident.

Ex. 2.11.10: Discrete sample spaces: suppose there are N cable cars in San Francisco, numbered sequentially from 1 to N . You see a cable car at random, it is numbered 203. You wish to estimate N .

- Assume your prior distribution on N is geometric with mean 100; that is,

$$p(N) = (1/100)(99/100)^{N-1}, \quad \text{for } N = 1, 2, \dots$$

What is your posterior distribution for N ?

- b What are the posterior mean and standard deviation of N ?
- c Choose a reasonable ‘noninformative’ prior distribution for N and give the resulting posterior distribution, mean, and standard deviation for N .

Answer:

Let $q = 99/100$. My first instinct was proposing an homogeneous sampling distribution for the number y of the observed cable car, $p(y|N) = 1/N$. But this conveys no information on the relation between y and N , and thus serves no good.

Ex. 2.11.11: Computing with a nonconjugate single parameter model: suppose y_1, \dots, y_5 are independent samples from a Cauchy distribution with unknown center θ and known scale 1: $p(y_i|\theta) \propto 1/(1 + (y_i - \theta)^2)$. Assume for simplicity that the prior distribution for θ is uniform on $[0, 100]$. Given the observations $(y_1, \dots, y_5) = (43, 44, 45, 46.5, 47.5)$:

- a Compute the unnormalized posterior density function, $p(\theta)p(y|\theta)$, on a grid of points $\theta = 0, \frac{1}{m}, \frac{2}{m}, \dots, 100$, for some large integer m . Using the grid approximation, compute and plot the normalized posterior density function, $p(\theta|y)$, as a function of θ .
- b Sample 1000 draws of θ from the posterior density and plot a histogram of the draws.
- c Use the 1000 samples of θ to obtain 1000 samples from the predictive distribution of a future observation, y_6 , and plot a histogram of the predictive draws.

Answer:

The following is an excerpt from the script “2.11.11.py” that does the math before plotting

```
import numpy as np

def normalized(arr):
    return arr/arr.sum()

m = 2**12 # This is unnecessarily big I guess

y = np.array([43, 44, 45, 46.5, 47.5])
theta = np.linspace(0, 100, m+1)
b_y, b_theta = np.meshgrid(y, theta) # "Big y, big theta"

theta_updf = 1
y_gvn_theta_updf = 1/np.prod(1+(b_y-b_theta)**2, axis=1)
# item a:
theta_gvn_y_pdf = normalized(theta_updf*y_gvn_theta_updf)

np.random.seed(42)

# item b:

def draw_samples(vals, pdf, size=1):
    cdf = np.cumsum(pdf)
    udraws = np.random.uniform(size=size)
    b_cdf, b_udraws = np.meshgrid(cdf, udraws)
    ids = np.argmax(b_cdf > b_udraws, axis=1)
    return vals[ids]

theta_gvn_y_samples = draw_samples(theta, theta_gvn_y_pdf, size=1000)

# item c:
y_pred_samples = np.random.standard_cauchy(size=1000)+theta_gvn_y_samples
```

From these data, Figure 6 is constructed.




Figure 6: Binned approximation to the posterior distribution for θ (a), histogram from such a distribution (b), and histogram for the posterior predictive distribution for $y_6 = \tilde{y}$ (c).

Ex. 2.11.12: Jeffreys' prior distributions: suppose $y|\theta \sim \text{Poisson}(\theta)$. Find Jeffreys' prior density for θ , and then find α and β for which the $\text{Gamma}(\alpha, \beta)$ density is a close match to Jeffreys' density.

Answer:

Jeffreys prescribes

$$p(\theta) \propto [J(\theta)]^{\frac{1}{2}},$$

where $J(\theta)$ is the Fisher Information for θ :

$$J(\theta) = \text{E} \left((\partial_{\theta} \log p(y|\theta))^2 \middle| \theta \right) = -\text{E} \left((\partial_{\theta})^2 \log p(y|\theta) \middle| \theta \right).$$

Direct calculation follows:

$$\begin{aligned} p(y|\theta) &= \text{Poisson}(\theta) \\ &= \frac{\theta^y e^{-\theta}}{y!}, \\ \log p(y|\theta) &= y \log \theta + f(y), \\ \partial_{\theta} \log p(y|\theta) &= \frac{y}{\theta}, \\ J(\theta) &= \frac{1}{\theta^2} \text{E}(y^2). \end{aligned}$$

It thus follows that Jeffreys' prior is the improper distribution $p(\theta) \propto 1/\theta$.

Since

$$\text{Gamma}(\theta|\alpha, \beta) \propto \theta^{\alpha-1} e^{-\beta\theta},$$

it follows that $\alpha \rightarrow 0^+$ and $\beta \rightarrow 0^+$ should approach this distribution.

Ex. 2.11.13: Discrete data: Table gives the number of fatal accidents and deaths on scheduled airline flights per year over a ten-year period. We use these data as a numerical example for fitting discrete data models.

- Assume that the numbers of fatal accidents in each year are independent with a $\text{Poisson}(\theta)$ distribution. Set a prior distribution for θ and determine the posterior distribution based on the data from 1976 through 1985. under this model, give a 95% predictive interval for the number of fatal accidents in 1986. You can use the normal approximation to the gamma and Poisson or compute using simulation.
- Assume that the numbers of fatal accidents in each year are independent with a $\text{Poisson}(\theta)$ distributions with a constant rate and an exposure in each year proportional to the number of passenger miles flown.

Year	Fatal accidents	Passenger deaths	Death rate
1976	24	734	0.19
1977	25	516	0.12
1978	31	754	0.15
1979	31	877	0.16
1980	22	814	0.14
1981	21	362	0.06
1982	26	764	0.15
1983	20	809	0.13
1984	16	223	0.03
1985	22	1066	0.15

Table 1: *Worldwide airline fatalities, 1976–1985. Death rate is passenger deaths per 100 million passenger miles. Source: Statistical Abstract of the United States.*

Set a prior distribution for θ and determine the posterior distribution based on the data for 1976 – 1985 . (Estimate the number of passenger miles flown in each year by dividing the appropriate columns of Table and ignoring round-off errors.) Give a 95% predictive interval for the number of fatal accidents in 1986 under the assumption that 8×10^{11} passenger miles are flown that year.

- c Repeat (a) above, replacing ‘fatal accidents’ with ‘passenger deaths.’
- d Repeat (b) above, replacing ‘fatal accidents’ with ‘passenger deaths.’
- e In which of the cases (a)–(d) above does the Poisson model seem more or less reasonable? Why? Discuss based on general principles, without sepecific reference to the numbers in Table

Incidentally, in 1986, there were 22 fatal accidents, 546 passenger deaths, and a death rate of 0.06 per 100 million miles flown. We return to this example in Exercises 3.12, 6.2, 6.3, and 8.14.